

Lianxiang Li

llxstupg@163.com | +86 175 1637 6491 | github.com/LI-PG1

Education

- **City University of Hong Kong** —MSc in Electronic Commerce (jointly offered by Dept. of Computer Science & Dept. of Information Systems), Fall 2026 – expected Oct 2027 / Feb 2028
- **Anhui University (Project 211)** —BSc in Applied Statistics, 2021 – 2025 | GPA 80/100

Experience

AI Development Engineer, VirtAItech (OrionX), Beijing

Mar – Aug 2026

- Delivered 30+ open LLMs (text, VLM, MoE, OCR, image/video) to production inference services; owned vLLM configuration, environment adaptation, and acceptance evaluation.
- Built reusable RAG and agent frameworks (embedding + rerank over Milvus, tool-calling orchestration); fine-tuned domain models with LoRA/QLoRA and DeepSpeed.
- Multimodal customer service for CITIC Securities: speech transcription (SenseVoice), document/chart field extraction (Qwen2.5-VL), finance knowledge base (73% auto-resolution, ~1.6s first response).

Selected Projects

Text2SQL Agent with Few-Shot Retrieval —Ctrip Data Platform

Jul – Aug 2026

- Schema linking and few-shot selection over historical query–SQL pairs; 400+ case golden evaluation set (indicator-QA 91%, SQL execution 88%).

Multi-Tool Agent for Equipment Maintenance —SAIC Motor, POC

Apr – Jun 2026

- Compared ReAct vs. Plan-then-Execute; MCP-style tool registration (7 tools); multi-step completion 74% → 86%.

Grounded RAG with Evidence Verification —Zhejiang Telecom

2026, during internship

- On-premise pipeline (4×A100); hybrid retrieval + cross-encoder rerank + LLM-as-Judge verification loop; top-5 hit rate +32%.

Research Interest

I am interested in Generative AI, LLMs, and agents. At work I built RAG and agent systems; I would now like to understand these models more deeply, not just deploy them. I am especially interested in how DSAIL applies generative models to scientific problems such as molecular design, and my statistics background makes me curious about the data side of that work.

Skills

Python | PyTorch, Transformers, DeepSpeed, PEFT (LoRA/QLoRA) | vLLM, Docker, FastAPI | LangChain, LangGraph, RAG, agent orchestration | Milvus, Chroma, Redis | probability, regression, multivariate analysis | IELTS 6.5

李联翔

邮箱: llxstug@163.com | 电话: +86 175 1637 6491 | GitHub: github.com/LI-PG1

教育背景

- 香港城市大学—电子商务理学硕士（计算机科学系与信息系统系合办），2026 秋—预计 2027.10 / 2028.02 毕业
- 安徽大学（211）—应用统计学学士，2021 —2025 | GPA 80/100

实习经历

- 大模型应用开发工程师，北京趋动科技（OrionX）2026.03 – 2026.08
- 交付 30+ 个开源大模型（文本/VLM/MoE/OCR/图像与视频）至生产推理服务，负责 vLLM 配置、环境适配与上线评估。
 - 搭建可复用 RAG 与 Agent 框架（Milvus 上的 Embedding+Rerank、工具调用编排），用 LoRA/QLoRA + DeepSpeed 微调垂域模型。
 - 参与中信建投多模态客服：语音转写（SenseVoice）、证件/图表字段抽取（Qwen2.5-VL）、金融知识库（自动解决率 73%，首响约 1.6s）。

代表项目

- Text2SQL 智能体 + Few-Shot 检索—携程数据中台2026.07 – 2026.08
- 设计 schema linking 与历史 query-SQL 对 few-shot 检索；构建 400+ 条黄金评估集（指标问答 91%，SQL 执行 88%）。
- 设备运维多工具智能体—上汽集团，POC2026.04 – 2026.06
- 对比 ReAct 与 Plan-then-Execute 编排；MCP 式工具注册（7 类工具）；多步任务完成率 74% → 86%。
- 带证据校验的落地 RAG —浙江电信2026（实习期间）
- 私有化部署（4×A100）；混合检索 + cross-encoder 精排 + LLM-as-Judge 校验闭环；top-5 命中率 +32%。

研究兴趣

我对生成式 AI、大语言模型和智能体感兴趣。工作里搭建过 RAG 和智能体系统；现在想更深入理解这些模型，而不仅仅是部署它们。我尤其关注 DSAIL 如何把生成模型应用到分子设计这类科学问题上，我的统计背景也让我对这项工作里的数据层面很好奇。

技能

Python | PyTorch、Transformers、DeepSpeed、PEFT（LoRA/QLoRA）| vLLM、Docker、FastAPI | LangChain、LangGraph、RAG、Agent 编排 | Milvus、Chroma、Redis | 概率论、回归、多元分析 | 雅思 6.5