

Lianxiang Li

llxstupg@163.com | +86 175 1637 6491 | github.com/LI-PG1

Education

- **City University of Hong Kong** —MSc in Electronic Commerce (jointly offered by Dept. of Computer Science & Dept. of Information Systems), Fall 2026 – expected Oct 2027 / Feb 2028
- **Anhui University (Project 211)** —BSc in Applied Statistics, 2021 – 2025 | GPA 80/100

Experience

AI Development Engineer, VirtAItech (OrionX), Beijing

Mar – Aug 2026

- Delivered 30+ open LLMs (text, VLM, MoE, OCR, image/video) to production inference services; owned vLLM configuration, environment adaptation, and acceptance evaluation.
- Built reusable RAG and agent frameworks (embedding + rerank over Milvus, tool-calling orchestration); fine-tuned domain models with LoRA/QLoRA and DeepSpeed.
- Multimodal customer service for CITIC Securities: speech transcription (SenseVoice), document/chart field extraction (Qwen2.5-VL), finance knowledge base (73% auto-resolution, ~1.6s first response).

Selected Projects

Text2SQL Agent with Few-Shot Retrieval —Ctrip Data Platform

Jul – Aug 2026

- Schema linking and few-shot selection over historical query–SQL pairs; 400+ case golden evaluation set (indicator-QA 91%, SQL execution 88%).

Multi-Tool Agent for Equipment Maintenance —SAIC Motor, POC

Apr – Jun 2026

- Compared ReAct vs. Plan-then-Execute; MCP-style tool registration (7 tools); multi-step completion 74% → 86%.

Grounded RAG with Evidence Verification —Zhejiang Telecom

2026, during internship

- On-premise pipeline (4×A100); hybrid retrieval + cross-encoder rerank + LLM-as-Judge verification loop; top-5 hit rate +32%.

Research Interest

I am interested in Generative AI, LLMs, and agents. At work I built RAG and agent systems; I would now like to understand these models more deeply, not just deploy them.

DemoDiff made me see that designing a model might itself be a learning task. A 0.7B model that outperforms much larger ones on 33 molecular design tasks left a strong impression on me. In Text2SQL I picked few-shot examples by hand. Some worked well, some did not, and I could not say why at the time. That is exactly what I want to understand.

I majored in statistics (regression, multivariate analysis), so data distribution and bias feel familiar. DSAIL's research depends a lot on data and evaluation, which is where I can help and where I want to learn. I hope to start as an RA and grow into research gradually.

Skills

Python | PyTorch, Transformers, DeepSpeed, PEFT (LoRA/QLoRA) | vLLM, Docker, FastAPI | LangChain, LangGraph, RAG, agent orchestration | Milvus, Chroma, Redis | probability, regression, multivariate analysis | IELTS 6.5

李联翔

邮箱: llxstupg@163.com | 电话: +86 175 1637 6491 | GitHub: github.com/LI-PG1

教育背景

- 香港城市大学—电子商务理学硕士（计算机科学系与信息系统系合办），2026 秋—预计 2027.10 / 2028.02 毕业
- 安徽大学（211）—应用统计学学士，2021 —2025 | GPA 80/100

实习经历

- | | |
|---|-------------------|
| 大模型应用开发工程师，北京趋动科技（OrionX） | 2026.03 – 2026.08 |
| <ul style="list-style-type: none">交付 30+ 个开源大模型（文本/VLM/MoE/OCR/图像与视频）至生产推理服务，负责 vLLM 配置、环境适配与上线评估。搭建可复用 RAG 与 Agent 框架（Milvus 上的 Embedding+Rerank、工具调用编排），用 LoRA/QLoRA + DeepSpeed 微调垂域模型。参与中信建投多模态客服：语音转写（SenseVoice）、证件/图表字段抽取（Qwen2.5-VL）、金融知识库（自动解决率 73%，首响约 1.6s）。 | |

代表项目

- | | |
|--|-------------------|
| Text2SQL 智能体 + Few-Shot 检索—携程数据中台 | 2026.07 – 2026.08 |
| <ul style="list-style-type: none">设计 schema linking 与历史 query-SQL 对 few-shot 检索；构建 400+ 条黄金评估集（指标问答 91%，SQL 执行 88%）。 | |
| 设备运维多工具智能体—上汽集团，POC | 2026.04 – 2026.06 |
| <ul style="list-style-type: none">对比 ReAct 与 Plan-then-Execute 编排；MCP 式工具注册（7 类工具）；多步任务完成率 74% → 86%。 | |
| 带证据校验的落地 RAG —浙江电信 | 2026（实习期间） |
| <ul style="list-style-type: none">私有化部署（4×A100）；混合检索 + cross-encoder 精排 + LLM-as-Judge 校验闭环；top-5 命中率 +32%。 | |

研究兴趣

我对生成式 AI、大语言模型和智能体感兴趣。工作里搭建过 RAG 和智能体系统；随着落地项目的增多，我想更深入理解这些模型，而不只是落地它们。

DemoDiff 让我看到一件事：‘设计模型’ 本身也许可以当作一个学习任务来做。0.7B 的小模型在 33 个分子设计任务上超过了大得多的模型，这个结果让我印象很深。我在 Text2SQL 里手动挑过 few-shot 例子，有些管用，有些不行，当时也说不出为什么。这正是我想弄明白的东西。

我本科是统计出身，学过回归、多元分析，对数据分布、偏差这类问题比较熟。DSAIL 的研究很看重数据和评估，这也是我能帮上忙、也愿意学的地方。希望先从 RA 做起，再慢慢深入了解研究。

技能

Python | PyTorch、Transformers、DeepSpeed、PEFT（LoRA/QLoRA）| vLLM、Docker、FastAPI | LangChain、LangGraph、RAG、Agent 编排 | Milvus、Chroma、Redis | 概率论、回归、多元分析 | 雅思 6.5